

CS34800

Information Systems

Object-Based Database Systems

Prof. Walid Aref
30 October 2017





Object-Oriented Databases



- Goal: Provide same benefits as object-oriented programming
 - Abstraction
 - Reuse
 - Natural data modeling
- A number of commercial systems
 - Gemstone (1986)
 - Informix, ObjectDB, O2, ...



Object-Oriented Databases



- Often programming language model
 - No separate query language
 - “Persistent Stores”
- But this gives up many of the advantages of Relational DB
 - Query optimization
 - Efficient concurrency control (we’ll discuss later)
 - Data independence



Solution: Object-Relational DB



- Incorporate key features into the relational model
 - User-defined data types
 - User-defined operations on those data types
- Postgres: Research project at Berkeley
 - Now available open source
- IBM bought Informix, Oracle included object-relational features
 - And almost nothing left of pure object-oriented DB



Object-Relational Data Models

- Extend the relational data model by including object orientation and constructs to deal with added data types.
- Allow attributes of tuples to have complex types, including non-atomic values such as nested relations.
- Preserve relational foundations, in particular the declarative access to data, while extending modeling power.
- Upward compatibility with existing relational languages.



Complex Data Types



- Goal: Intuitive modeling of complex data
 - Abstraction
- Basic idea: Non-atomic domains
 - Cell can contain something other than “atomic” (indivisible) value
 - Example of non-atomic domain: set of integers, or set of tuples
- What part of the relational model does this violate?
 - A. Everything represented as a relation (table)
 - B. First normal form
 - C. Relational algebra
 - D. Declarative query language
- “Standardized” in SQL:1999
 - But most commercial systems vary from standard



Complex Types and SQL

- Extensions introduced in SQL:1999 to support complex types:
 - Collection and large object types
 - ▶ Nested relations are an example of collection types
 - Structured types
 - ▶ Nested record structures like composite attributes
 - Inheritance
 - Object orientation
 - ▶ Including object identifiers and references
- Commercial databases may vary from the standard
 - But some features are present in each of the major commercial database systems
 - ▶ Read the manual of your database system to see what it supports



Array and Multiset Types in SQL

- Example of array and multiset declaration:

```
create type Publisher as  
    (name          varchar(20),  
     branch       varchar(20));  
create type Book as  
    (title         varchar(20),  
     author_array varchar(20) array [10],  
     pub_date      date,  
     publisher     Publisher,  
     keyword-set   varchar(20) multiset);  
create table books of Book;
```




Creation of Collection Values

- Array construction
array ['Silberschatz', `Korth`, `Sudarshan']
- Multisets
multiset ['computer', 'database', 'SQL']
- To create a tuple of the type defined by the books relation:
('Compilers', **array**[`Smith`, `Jones`],
new Publisher (`McGraw-Hill`, `New York`),
multiset [`parsing`, `analysis'])
- To insert the preceding tuple into the relation books
insert into *books*
values
('Compilers', **array**[`Smith`, `Jones`],
new Publisher (`McGraw-Hill`, `New York`),
multiset [`parsing`, `analysis']);



Example of a Nested Relation

- Example: library information system
- Each book has
 - title,
 - a list (array) of authors,
 - Publisher, with subfields *name* and *branch*, and
 - a set of keywords
- Non-1NF relation *books*
 - *Idea: Model as cells that contain relations*

<i>title</i>	<i>author_array</i>	<i>publisher</i>	<i>keyword_set</i>
		(<i>name, branch</i>)	
Compilers	[Smith, Jones]	(McGraw-Hill, NewYork)	{parsing, analysis}
Networks	[Jones, Frick]	(Oxford, London)	{Internet, Web}



4NF Decomposition of Nested Relation

- Suppose for simplicity that title uniquely identifies a book
 - In real world ISBN is a unique identifier
- Decompose *books* into 4NF using the schemas:
 - $(title, author, position)$
 - $(title, keyword)$
 - $(title, pub_name, pub_branch)$
- 4NF design requires users to include joins in their queries.

<i>title</i>	<i>author</i>	<i>position</i>
Compilers	Smith	1
Compilers	Jones	2
Networks	Jones	1
Networks	Frick	2

authors

<i>title</i>	<i>keyword</i>
Compilers	parsing
Compilers	analysis
Networks	Internet
Networks	Web

keywords

<i>title</i>	<i>pub_name</i>	<i>pub_branch</i>
Compilers	McGraw-Hill	New York
Networks	Oxford	London

books4



Querying Collection-Valued Attributes

- To find all books that have the word “database” as a keyword,
select *title*
from *books*
where ‘database’ **in** (**unnest**(*keyword-set*))
- We can access individual elements of an array by using indices
 - E.g.: If we know that a particular book has three authors, we could write:
select *author_array*[1], *author_array*[2], *author_array*[3]
from *books*
where *title* = ‘Database System Concepts’
- To get a relation containing pairs of the form “title, author_name” for each book and each author of the book
select *B.title*, *A.author*
from *books* **as** *B*, **unnest** (*B.author_array*) **as** *A* (*author*)
- To retain ordering information we add a **with ordinality** clause
select *B.title*, *A.author*, *A.position*
from *books* **as** *B*, **unnest** (*B.author_array*) **with ordinality** **as**
A (*author*, *position*)



Unnesting

- The transformation of a nested relation into a form with fewer (or no) relation-valued attributes is called **unnesting**.
- E.g.
select *title*, *A* **as** *author*, *publisher.name* **as** *pub_name*,
publisher.branch **as** *pub_branch*, *K.keyword*
from *books* **as** *B*, **unnest**(*B.author_array*) **as** *A* (*author*),
unnest (*B.keyword_set*) **as** *K* (*keyword*)
- Result relation *flat_books*

<i>title</i>	<i>author</i>	<i>pub_name</i>	<i>pub_branch</i>	<i>keyword</i>
Compilers	Smith	McGraw-Hill	New York	parsing
Compilers	Jones	McGraw-Hill	New York	parsing
Compilers	Smith	McGraw-Hill	New York	analysis
Compilers	Jones	McGraw-Hill	New York	analysis
Networks	Jones	Oxford	London	Internet
Networks	Frick	Oxford	London	Internet
Networks	Jones	Oxford	London	Web
Networks	Frick	Oxford	London	Web



Nesting

- **Nesting** is the opposite of unnesting, creating a collection-valued attribute
- Nesting can be done in a manner similar to aggregation, but using the function **collect()** in place of an aggregation operation, to create a multiset
- To nest the *flat_books* relation on the attribute *keyword*:

```
select title, author, Publisher (pub_name, pub_branch) as publisher,  
       collect (keyword) as keyword_set  
from flat_books  
groupby title, author, publisher
```

- To nest on both authors and keywords:

```
select title, collect (author) as author_set,  
       Publisher (pub_name, pub_branch) as publisher,  
       collect (keyword) as keyword_set  
from flat_books  
group by title, publisher
```



Nesting (Cont.)

- Another approach to creating nested relations is to use subqueries in the **select** clause, starting from the 4NF relation *books4*

```
select title,  
      array (select author  
              from authors as A  
              where A.title = B.title  
              order by A.position) as author_array,  
      Publisher (pub-name, pub-branch) as publisher,  
      multiset (select keyword  
                  from keywords as K  
                  where K.title = B.title) as keyword_set  
from books4 as B
```



Structured Types in SQL

- **Structured types** (a.k.a. **user-defined types**) can be declared and used in SQL

create type *Name* **as**

```
(firstname      varchar(20),  
  lastname      varchar(20))  
final
```

create type *Address* **as**

```
(street         varchar(20),  
  city          varchar(20),  
  zipcode       varchar(20))  
not final
```

- Note: **final** and **not final** indicate whether subtypes can be created
- Structured types can be used to create tables with composite attributes

```
create table person (  
  name      Name,  
  address   Address,  
  dateOfBirth date)
```
- Dot notation used to reference components: *name.firstname*



Structured Types (cont.)

■ User-defined row types

```
create type PersonType as (  
    name Name,  
    address Address,  
    dateOfBirth date)  
not final
```

■ Can then create a table whose rows are a user-defined type

```
create table customer of CustomerType
```

■ Alternative using **unnamed row types**.

```
create table person_r(  
    name    row(firstname varchar(20),  
                lastname varchar(20)),  
    address row(street    varchar(20),  
                city       varchar(20),  
                zipcode   varchar(20)),  
    dateOfBirth date)
```



Methods

- Can add a method declaration with a structured type.

method *ageOnDate* (*onDate* **date**)

returns **interval year**

- Method body is given separately.

create instance method *ageOnDate* (*onDate* **date**)

returns **interval year**

for *CustomerType*

begin

return *onDate* - **self.dateOfBirth**;

end

- We can now find the age of each customer:

select *name.lastname*, *ageOnDate* (**current_date**)

from *customer*



Constructor Functions

- **Constructor functions** are used to create values of structured types
- E.g.
create function *Name*(*firstname* **varchar**(20), *lastname* **varchar**(20))
returns *Name*
begin
 set *self.firstname* = *firstname*;
 set *self.lastname* = *lastname*;
end
- To create a value of type *Name*, we use
 new *Name*('John', 'Smith')
- Normally used in insert statements
insert into *Person* **values**
 (**new** *Name*('John', 'Smith'),
 new *Address*('20 Main St', 'New York', '11001'),
 date '1960-8-22');



Type Inheritance

- Suppose that we have the following type definition for people:

```
create type Person  
  (name varchar(20),  
   address varchar(20))
```

- Using inheritance to define the student and teacher types

```
create type Student  
under Person  
  (degree varchar(20),  
   department varchar(20))
```

```
create type Teacher  
under Person  
  (salary integer,  
   department varchar(20))
```

- Subtypes can redefine methods by using **overriding method** in place of **method** in the method declaration



Multiple Type Inheritance

- SQL:1999 and SQL:2003 do not support multiple inheritance
- If our type system supports multiple inheritance, we can define a type for teaching assistant as follows:

```
create type Teaching Assistant  
under Student, Teacher
```

- To avoid a conflict between the two occurrences of *department* we can rename them

```
create type Teaching Assistant  
under  
  Student with (department as student_dept),  
  Teacher with (department as teacher_dept)
```

- Each value must have a **most-specific type**



Table Inheritance

- Tables created from subtypes can further be specified as **subtables**
- E.g. **create table** *people* **of** *Person*;
 create table *students* **of** *Student* **under** *people*;
 create table *teachers* **of** *Teacher* **under** *people*;
- Tuples added to a subtable are automatically visible to queries on the supertable
 - E.g. query on *people* also sees *students* and *teachers*.
 - Similarly updates/deletes on *people* also result in updates/deletes on subtables
 - To override this behaviour, use “**only** *people*” in query
- Conceptually, multiple inheritance is possible with tables
 - e.g. *teaching_assistants* under *students* and *teachers*
 - *But is not supported in SQL currently*
 - ▶ So we cannot create a person (tuple in *people*) who is both a student and a teacher



Consistency Requirements for Subtables

- Consistency requirements on subtables and supertables.
 - Each tuple of the supertable (e.g. *people*) can correspond to at most one tuple in each of the subtables (e.g. *students* and *teachers*)
 - Additional constraint in SQL:1999:

All tuples corresponding to each other (that is, with the same values for inherited attributes) must be derived from one tuple (inserted into one table).

 - ▶ That is, each entity must have a most specific type
 - ▶ We cannot have a tuple in *people* corresponding to a tuple each in *students* and *teachers*



Object-Identity and Reference Types

- Define a type *Department* with a field *name* and a field *head* which is a reference to the type *Person*, with table *people* as scope:

```
create type Department (  
    name varchar (20),  
    head ref (Person) scope people)
```

- We can then create a table *departments* as follows

```
create table departments of Department
```

- We can omit the declaration **scope** *people* from the type declaration and instead make an addition to the **create table** statement:

```
create table departments of Department  
    (head with options scope people)
```

- Referenced table must have an attribute that stores the identifier, called the **self-referential attribute**

```
create table people of Person  
    ref is person_id system generated;
```




Initializing Reference-Typed Values

- To create a tuple with a reference value, we can first create the tuple with a null reference and then set the reference separately:

```
insert into departments  
  values (`CS`, null)  
update departments  
  set head = (select p.person_id  
                from people as p  
                where name = `John`)  
  where name = `CS`
```



User Generated Identifiers

- The type of the object-identifier must be specified as part of the type definition of the referenced table, and
- The table definition must specify that the reference is user generated

```
create type Person  
  (name varchar(20)  
   address varchar(20))  
  ref using varchar(20)  
create table people of Person  
  ref is person_id user generated
```

- When creating a tuple, we must provide a unique value for the identifier:

```
insert into people (person_id, name, address) values  
  ('01284567', 'John', '23 Coyote Run')
```

- We can then use the identifier value when inserting a tuple into *departments*

- Avoids need for a separate query to retrieve the identifier:

```
insert into departments  
values('CS', '02184567')
```



User Generated Identifiers (Cont.)

- Can use an existing primary key value as the identifier:

```
create type Person  
    (name varchar (20) primary key,  
     address varchar(20))  
ref from (name)  
create table people of Person  
    ref is person_id derived
```

- When inserting a tuple for *departments*, we can then use

```
insert into departments  
    values(`CS`,`John`)
```



Path Expressions

- Find the names and addresses of the heads of all departments:
select *head* → *name*, *head* → *address*
from *departments*
- An expression such as “*head* → *name*” is called a **path expression**
- Path expressions help avoid explicit joins
 - If department head were not a reference, a join of *departments* with *people* would be required to get at the address
 - Makes expressing the query much easier for the user



Persistent Programming Languages

- Languages extended with constructs to handle persistent data
- Programmer can manipulate persistent data directly
 - no need to fetch it into memory and store it back to disk (unlike embedded SQL)
- Persistent objects:
 - **Persistence by class** - explicit declaration of persistence
 - **Persistence by creation** - special syntax to create persistent objects
 - **Persistence by marking** - make objects persistent after creation
 - **Persistence by reachability** - object is persistent if it is declared explicitly to be so or is reachable from a persistent object



Object Identity and Pointers

- Degrees of permanence of object identity
 - **Intraprocedure**: only during execution of a single procedure
 - **Intraprogram**: only during execution of a single program or query
 - **Interprogram**: across program executions, but not if data-storage format on disk changes
 - **Persistent**: interprogram, plus persistent across data reorganizations
- Persistent versions of C++ and Java have been implemented
 - C++
 - ▶ ODMG C++
 - ▶ ObjectStore
 - Java
 - ▶ Java Database Objects (JDO)



Persistent C++ Systems

- Extensions of C++ language to support persistent storage of objects
- Several proposals, ODMG standard proposed, but not much action of late
 - **persistent pointers**: e.g. `d_Ref<T>`
 - **creation of persistent objects**: e.g. `new (db) T()`
 - **Class extents**: access to all persistent objects of a particular class
 - **Relationships**: Represented by pointers stored in related objects
 - ▶ Issue: consistency of pointers
 - ▶ Solution: extension to type system to automatically maintain back-references
 - **Iterator interface**
 - **Transactions**
 - **Updates**: `mark_modified()` function to tell system that a persistent object that was fetched into memory has been updated
 - **Query language**



Persistent Java Systems

- Standard for adding persistence to Java : **Java Database Objects (JDO)**
 - Persistence by reachability
 - Byte code enhancement
 - ▶ Classes separately declared as persistent
 - ▶ Byte code modifier program modifies class byte code to support persistence
 - E.g. Fetch object on demand
 - Mark modified objects to be written back to database
 - Database mapping
 - ▶ Allows objects to be stored in a relational database
 - Class extents
 - Single reference type
 - ▶ no difference between in-memory pointer and persistent pointer
 - ▶ Implementation technique based on **hollow objects** (a.k.a. **pointer swizzling**)



Object-Relational Mapping

- **Object-Relational Mapping (ORM)** systems built on top of traditional relational databases
- Implementor provides a mapping from objects to relations
 - Objects are purely transient, no permanent object identity
- Objects can be retrieved from database
 - System uses mapping to fetch relevant data from relations and construct objects
 - Updated objects are stored back in database by generating corresponding update/insert/delete statements
- The **Hibernate** ORM system is widely used
 - described in Section 9.4.2
 - Provides API to start/end transactions, fetch objects, etc
 - Provides query language operating directly on object model
 - ▶ queries translated to SQL
- Limitations: overheads, especially for bulk updates



Comparison of O-O and O-R Databases

■ Relational systems

- simple data types, powerful query languages, high protection.

■ Persistent-programming-language-based OODBs

- complex data types, integration with programming language, high performance.

■ Object-relational systems

- complex data types, powerful query languages, high protection.

■ Object-relational mapping systems

- complex data types integrated with programming language, but built as a layer on top of a relational database system

■ Note: Many real systems blur these boundaries

- E.g. persistent programming language built as a wrapper on a relational database offers first two benefits, but may have poor performance.